

小文件解决方案

小文件解决方案

1. 概述

小文件是指文件size小于HDFS上block大小的文件。这样的文件会给hadoop的扩展性和性能带来严重问题。首先，在HDFS中，任何block，文件或者目录在内存中均以对象的形式存储，每个对象约占150byte，如果有1百万个小文件，每个文件占用一个block，则namenode大约需要2G空间。如果存储1亿个文件，则namenode需要20G空间，这样namenode内存容量严重制约了集群的扩展。其次，访问大量小文件速度远远小于访问几个大文件。HDFS最初是为流式访问大文件开发的，如果访问大量小文件，需要不断的从一个datanode跳到另一个datanode，严重影响性能。每一个小文件要占用一个slot，而task启动将耗费大量时间甚至大部分时间都耗费在启动task和释放task上。

2. Hadoop的解决方案

对于小文件问题，Hadoop提供了几个解决方案，分别为：Hadoop Archive，Sequence file和CombineFileInputFormat。

2.1. Hadoop Archive

Hadoop Archive或者HAR，是一个高效地将小文件放入HDFS块中的文件存档工具，它能够将多个小文件打包成一个HAR文件，这样在减少namenode内存使用的同时，仍然允许对文件进行透明的访问。

2.1.1. 创建har文件

语法：

```
hadoop archive -archiveName <name> <src> <dest>
```

事例：

对某个目录/ceshi下的所有小文件存档成/ceshi/out1:

```
hadoop archive -archiveName ceshi.har -p hdfs://master:9000/ceshi2 hdfs://master:9000/ceshi
```

2.1.2. 查看har文件

语法：

Hadoop Archives的URI是har://scheme-hostname:port/archivepath/fileinarchive。

如果没提供scheme-hostname，它会使用默认的文件系统。这种情况下URI是这种形式har:///archivepath/fileinarchive。

例如：

```
hadoop fs -ls har://hdfs-master:9000/ceshi/ceshi.har
```

```
hadoop fs -ls har:///ceshi/ceshi.har
```

```
[hadoop@master ~]$ hadoop fs -ls har:///ceshi/ceshi.har
```

```
Warning: $HADOOP_HOME is deprecated.
```

Found 3 items

```
-rw-r--r-- 2 hadoop supergroup 22930015 2015-06-08 19:43 /ceshi/ceshi.har/partition.txt
```

```
-rw-r--r-- 2 hadoop supergroup 25 2015-06-09 12:03 /ceshi/ceshi.har/mini.txt
```

```
-rw-r--r-- 2 hadoop supergroup 14340 2015-06-10 06:32 /ceshi/ceshi.har/access.log
```

```
[hadoop@master ~]$ hadoop fs -cat har:///ceshi/ceshi.har/mini.txt
```

```
Warning: $HADOOP_HOME is deprecated.
```

```
hello world
```

2.1.3. mapreduce操作har

wordcount

```
package com.itcast.wordcount;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
```

```

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.util.GenericOptionsParser;

public class WordCount {

    public static class TokenizerMapper
    extends Mapper<Object, Text, Text, IntWritable>{

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context
        ) throws IOException, InterruptedException {
            //value: 一行一行数据，对数据做切分
            String[] split = value.toString().split(" ");
            //来源url
            String url = split[6];
            //输出
            word.set(url);
            context.write(word, one);
        }
    }

    public static class IntSumReducer
    extends Reducer<Text,IntWritable,Text,IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
        Context context
        ) throws IOException, InterruptedException {
            //循环相加每一个int类型的值
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            //设置url
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws Exception {
        Configuration conf = new Configuration();
        conf.set("df.default.name", "hdfs://master:9000/");//设置hdfs的默认路径
        conf.set("hadoop.job.ugi", "hadoop,hadoop");//用户名，组
        conf.set("mapred.job.tracker", "master:9001");//设置jobtracker在哪
        String[] otherArgs = { "har://hdfs-master:9000/ceshi/ceshi.har",//
        输入文件路径
        "hdfs://master:9000/ceshi/out9/" };//输入输出路径，可以设置，也可以通过main方法传进来
        // String[] otherArgs = new GenericOptionsParser(conf, args).getRemainingArgs();//输入和输出参数
        if (otherArgs.length != 2) {
            System.err.println("Usage: wordcount <in> <out>");
            System.exit(2);
        }
        Job job = new Job(conf, "word count");
        job.setJarByClass(WordCount.class);
        job.setMapperClass(TokenizerMapper.class);
        job.setCombinerClass(IntSumReducer.class);
        job.setReducerClass(IntSumReducer.class);
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.class);
        FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
        FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}

```

2.1.4. 注意事项

使用HAR时需要两点，第一，对小文件进行存档后，原文件并不会自动被删除，需要用户自己删除；第二，创建HAR文件的过程实际上是在运行一个mapreduce作业，因而需要有一个hadoop集群运行此命令。

此外，HAR还有一些缺陷：第一，一旦创建，Archives便不可改变。要增加或删除里面的文件，必须重新创建归档文件。第二，要归档的文件名中不能有空格，否则会抛出异常

2.2. Sequence file

sequence file由一系列的二进制key/value组成，如果为key小文件名，value为文件内容，则可以将大批小文件合并成一个大文件。

Hadoop中提供了SequenceFile，包括Writer，Reader和SequenceFileSorter类进行写，读和排序操作。

2.2.1. 序列化文件读写

```
package com.itcast.userdefined.sequence.makedata;

import java.io.IOException;
import java.io.InputStream;
import java.net.URI;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.FileStatus;
import org.apache.hadoop.fs.FileSystem;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IOUtils;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.SequenceFile;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.Writable;
import org.apache.hadoop.util.ReflectionUtils;
import org.junit.Test;

public class SeqComSmallFile {
    @Test
    public void testWriter(){

        Configuration conf = new Configuration();//创建配置信息
        conf.set("fs.default.name", "hdfs://master:9000");//不写走默认
        conf.set("hadoop.job.ugi", "hadoop,hadoop");//如果不写系统将按照默认的用户进行操作
        String url = "hdfs://master:9000/ceshi";//读取的文件
        Path path = new Path("hdfs://master:9000/ceshi2/test.seq");//序列化文件名
        FileSystem fs = null;
        InputStream is = null;
        SequenceFile.Writer w = null;//创建writer流
        Text k = new Text();//key，相当于string
        Text v = new Text();//value，相当于String
        try {
            //代表一个统一资源标识符(URI)
            fs = FileSystem.get(URI.create(url), conf);//创建filesystem
            w = SequenceFile.createWriter(fs, conf, path, k.getClass(), v.getClass());//创建writer
            for (FileStatus a : fs.listStatus(new Path(url))) {
                is = fs.open(a.getPath());//通过filesystem打开文件读取流
                k.set(a.getPath().getName());
                byte buf[] = new byte[1024*64];
                int bytesRead = is.read(buf);
                System.out.println();
                v.set(new String(buf));
                w.append(k, v);
                IOUtils.copyBytes(is, System.out, 1024, false);
            }

        } catch (Exception e) {
            e.printStackTrace();
        } finally {
            try {
                IOUtils.closeStream(is);
                fs.close();
            } catch (IOException e1) {
                fs=null;
                e1.printStackTrace();
            }
        }

    }

    @Test
    public void TestReader(){

        Configuration conf = new Configuration();//创建配置信息
        conf.set("fs.default.name", "hdfs://master:9000");//默认路径
        conf.set("hadoop.job.ugi", "hadoop,hadoop");//用户和组的信息
        String uriin = "hdfs://master:9000/ceshi2/";//文件位置
        FileSystem fs = null;
        try {
            fs = FileSystem.get(URI.create(uriin), conf);
        } catch (IOException e2) {
            e2.printStackTrace();
        } //创建filesystem
        Path path = new Path("hdfs://master:9000/ceshi/out1/part-r-00000");//序列化文件的路径
        SequenceFile.Reader r = null;
```

```

try {
    r = new SequenceFile.Reader(fs, path, conf); //通过序列化文件类
    SequenceFile创建readerliu
    Writable k = (Writable) ReflectionUtils.newInstance(r.getKeyClas
    s(), conf); //创建key
    Writable v = (Writable) ReflectionUtils.newInstance(r.getValueCl
    ass(), conf); //创建value
    while (r.next(k, v)) { //顺序读取
        System.out.println(k + "\n" + v);
    }
} catch (Exception e) {
    e.printStackTrace();
} finally {
    IOUtils.closeStream(r);
}
try {
    fs.close();
} catch (IOException e1) {
    fs=null;
    e1.printStackTrace();
}
}
}
}

```

2.2.2. mapreduce输入序列化文件

```

package com.itcast.userdefined.sequence.input;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.input.SequenceFileInp
utFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFor
mat;

/**
 * 处理序列化文件
 * @author wilson
 */
public class SequenceFileInputFormatDemo {
    public static class TokenizerMapper extends
    Mapper<Text, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();
        //key为序列化文件的key，也就是文件的名字，value是序列化文
        件的value，也就是文件的内容
        public void map(Text key, Text value, Context context)
        throws IOException, InterruptedException {
            //用'\n'将value分割成每一行
            String[] split = value.toString().split("\n");
            for (String string : split) {
                //用" "将每一行分割
                String[] split2 = string.split(" ");
                //有空字符串存在,判空
                if (split2.length>=6) {
                    //来源url
                    String url = split2[6];
                    //设置、输出
                    word.set(url);
                    context.write(word, one);
                } else {
                    word.set(string);
                    context.write(word, one);
                }
            }
        }

    }

    public static class IntSumReducer extends
    Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();
    }
}

```

```

public void reduce(Text key, Iterable<IntWritable> values,
Context context) throws IOException, InterruptedException {
// 循环相加每一个int类型的值
int sum = 0;
for (IntWritable val : values) {
sum += val.get();
}
// 设置url
result.set(sum);
context.write(key, result);
}
}

public static void main(String[] args) throws Exception {
Configuration conf = new Configuration();
conf.set("fs.default.name", "hdfs://master:9000/");// 设置hdfs的
默认路径
conf.set("hadoop.job.ugi", "hadoop,hadoop");// 用户名, 组
conf.set("mapred.job.tracker", "master:9001");// 设置jobtracker
在哪
String[] otherArgs = { "hdfs://master:9000/ceshi2/test.seq",// 输
入文件路径
"hdfs://master:9000/ceshi/out1/" };// 输入输出路径, 可以设置,
也可以通过main方法传进来
// String[] otherArgs = new GenericOptionsParser(conf,
// args).getRemainingArgs();//输入和输出参数
Job job = new Job(conf, "sequence file input demo");
job.setJarByClass(SequenceFileInputFormatDemo.class);

job.setMapperClass(TokenizerMapper.class);
job.setReducerClass(IntSumReducer.class);
job.setNumReduceTasks(1);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);
job.setInputFormatClass(SequenceFileInputFormat.class);
FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

2.2.3. mapreduce输出序列化文件

```

package com.itcast.userdefined.sequence.output;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.SequenceFile.CompressionType;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.io.compress.DefaultCodec;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
import org.apache.hadoop.mapreduce.lib.output.SequenceFileOutputFormat;

public class SequenceOutputFormat {
    public static class TokenizerMapper extends
Mapper<Object, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(Object key, Text value, Context context)
throws IOException, InterruptedException {
// value: 一行一行数据, 对数据做切分
String[] split = value.toString().split(" ");
// 来源url
String url = split[6];
// 输出
word.set(url);
context.write(word, one);
}
}
}

```

```

public static class IntSumReducer extends
Reducer<Text, IntWritable, Text, IntWritable> {
private IntWritable result = new IntWritable();

public void reduce(Text key, Iterable<IntWritable> values,
Context context) throws IOException, InterruptedException {
// 循环相加每一个int类型的值
int sum = 0;
for (IntWritable val : values) {
sum += val.get();
}
// 设置url
result.set(sum);
context.write(key, result);
}
}

public static void main(String[] args) throws Exception {
Configuration conf = new Configuration();
conf.set("fs.default.name", "hdfs://master:9000/");// 设置hdfs的
默认路径
conf.set("hadoop.job.ugi", "hadoop,hadoop");// 用户名, 组
conf.set("mapred.job.tracker", "master:9001");// 设置jobtracker
在哪
String[] otherArgs = { "hdfs://master:9000/ceshi/access.log",// 输
入文件路径
"hdfs://master:9000/ceshi/out3/" };// 输入输出路径, 可以设置,
也可以通过main方法传进来
// String[] otherArgs = new GenericOptionsParser(conf,
// args).getRemainingArgs();//输入和输出参数
if (otherArgs.length != 2) {
System.err.println("Usage: wordcount <in> <out>");
System.exit(2);
}
Job job = new Job(conf, "sequence file output demo ");
job.setJarByClass(SequenceOutputFormat.class);
job.setMapperClass(TokenizerMapper.class);
job.setReducerClass(IntSumReducer.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
// 设置输出类
job.setOutputFormatClass(SequenceFileOutputFormat.class);
// 设置sequencefile的格式, 对于sequencefile的输出格式, 有多种
组合方式.
// 从下面的模式中选择一种, 并将其余的注释掉
// 组合方式1: 不压缩模式
//SequenceFileOutputFormat.setOutputCompressionType(job,C
ompressionType.NONE);

// 组合方式2: record压缩模式(只压缩值),
//并指定采用的压缩方式 : 默认:gzip压缩等

// SequenceFileOutputFormat.setOutputCompressionType(job,
// CompressionType.RECORD);
// SequenceFileOutputFormat.setOutputCompressorClass(job,
// DefaultCodec.class);

// 组合方式3: block压缩模式(压缩序列记录到一个块),
//并指定采用的压缩方式 : 默认、gzip压缩等

// SequenceFileOutputFormat.setOutputCompressionType(job,
// CompressionType.BLOCK);
// SequenceFileOutputFormat.setOutputCompressorClass(job,
// DefaultCodec.class);

FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}

```

2.3. CombineFileInputFormat

2.3.1. 介绍

CombineFileInputFormat是一种新的inputformat, 用于将多个文件合并成一个单独的split, 另外, 它会考虑数据的存储位置。

CombineFileInputFormat的大致原理是, 他会将输入多个数据文件(小文件)的元数据全部包装到CombineFileSplit类里面。也就是说, 因为小文件的情况下, 在HDFS中都是单Block的文件, 即一个文件一个Block, 一个CombineFileSplit包含了一组文件Block, 包括每个文件的起始偏移(offset), 长度(length), Block位置(localtions)等元数据。如果想要处理一个CombineFileSplit, 很容易想

到，对其包含的每个InputSplit（实际上这里面没有这个，你需要读取一个小文件块的时候，需要构造一个FileInputSplit对象）。

在执行MapReduce任务的时候，需要读取文件的文本行（简单一点是文本行，也可能是其他格式数据）。那么对于CombineFileSplit来说，你需要处理其包含的小文件Block，就要对应设置一个RecordReader，才能正确读取文件数据内容。通常情况下，我们有一批小文件，格式通常是相同的，只需要在为CombineFileSplit实现一个RecordReader的时候，内置另一个用来读取小文件Block的RecordReader，这样就能保证读取CombineFileSplit内部聚积的小文件。

2.3.2. 代码实现

mapreduce代码

```
package com.itcast.userdefined.CombineInputFormat;

import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class CombineSmallfiles {
    public static class TokenizerMapper extends
        Mapper<LongWritable, Text, Text, IntWritable> {

        private final static IntWritable one = new IntWritable(1);
        private Text word = new Text();

        public void map(LongWritable key, Text value, Context context)
            throws IOException, InterruptedException {
            // value: 一行一行数据，对数据做切分
            String[] split = value.toString().split(" ");
            // 来源url
            String url = split[6];
            // 输出
            word.set(url);
            context.write(word, one);
        }
    }

    public static class IntSumReducer extends
        Reducer<Text, IntWritable, Text, IntWritable> {
        private IntWritable result = new IntWritable();

        public void reduce(Text key, Iterable<IntWritable> values,
            Context context) throws IOException, InterruptedException {
            // 循环相加每一个int类型的值
            int sum = 0;
            for (IntWritable val : values) {
                sum += val.get();
            }
            // 设置url
            result.set(sum);
            context.write(key, result);
        }
    }

    public static void main(String[] args) throws IOException,
        ClassNotFoundException, InterruptedException {

        Configuration conf = new Configuration();
        conf.set("df.default.name", "hdfs://master:9000/");// 设置hdfs的
        默认路径
        conf.set("hadoop.job.ugi", "hadoop,hadoop");// 用户名，组
        conf.set("mapred.job.tracker", "master:9001");// 设置jobtracker
        在哪
        /*
        * 在同一个节点上的Blocks合并，超过maxSplitSize就生成新分片
        。如果没有指定，则只汇总本节点Block，暂不分片。
        */
        conf.setLong("mapreduce.input.fileinputformat.split.maxsize",
            64 * 1024 * 1024);
        /*
        * 把maxsize中处理剩余的Block，进行合并，如果超过minSizeNode，
        那么全部作为一个分片。
        */
        conf.setLong("mapreduce.input.fileinputformat.split.minsize.per.
            node",
```

```

1024);

// conf.setInt("mapred.min.split.size", 1);
// conf.setLong("mapred.max.split.size", 26214400); // 25m
String[] otherArgs = { "hdfs://master:9000/ceshi", // 输入文件路径
"hdfs://master:9000/ceshi/out1" }; // 输入输出路径，可以设置，
也可以通过main方法传进来
// String[] otherArgs = new GenericOptionsParser(conf,
// args).getRemainingArgs(); // 输入和输出参数
if (otherArgs.length != 2) {
System.err.println("Usage: wordcount <in> <out>");
System.exit(2);
}
Job job = new Job(conf, "combine smallfiles");
job.setJarByClass(CombineSmallfiles.class);
job.setMapperClass(TokenizerMapper.class);
job.setReducerClass(IntSumReducer.class);
job.setMapOutputKeyClass(Text.class);
job.setMapOutputValueClass(IntWritable.class);
job.setOutputKeyClass(Text.class);
job.setOutputValueClass(IntWritable.class);
job.setInputFormatClass(CombineSmallfileInputFormat.class);
FileInputFormat.addInputPath(job, new Path(otherArgs[0]));
FileOutputFormat.setOutputPath(job, new Path(otherArgs[1]));
System.exit(job.waitForCompletion(true) ? 0 : 1);

}

}

```

inputformat代码

```

package com.itcast.userdefined.CombineInputFormat;

import java.io.IOException;

import org.apache.hadoop.io.BytesWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.InputSplit;
import org.apache.hadoop.mapreduce.RecordReader;
import org.apache.hadoop.mapreduce.TaskAttemptContext;
import org.apache.hadoop.mapreduce.lib.input.CombineFileInput
tFormat;
import org.apache.hadoop.mapreduce.lib.input.CombineFileRec
ordReader;
import org.apache.hadoop.mapreduce.lib.input.CombineFileSplit
;
/**
 * 继承CombineFileInputFormat，采用combineFileSplit，实现自
 * 己的recordReader
 * @author wilson
 */
public class CombineSmallfileInputFormat extends CombineFileI
nputFormat<LongWritable, Text> {

    @Override
    public RecordReader<LongWritable, Text> createRecordReade
r(InputSplit split, TaskAttemptContext context) throws IOExceptio
n {

        CombineFileRecordReader<LongWritable, Text> recordReader
        = new CombineFileRecordReader<LongWritable, Text>((Combi
        neFileSplit) split, context, CombineSmallfileRecordReader.class)
        ;
        try {
            recordReader.initialize((CombineFileSplit) split, context);
        } catch (InterruptedException e) {
            new RuntimeException("Error to initialize CombineSmallfileRecor
            dReader.");
        }
        return recordReader;
    }

}

```

recordreader代码

```

package com.itcast.userdefined.CombineInputFormat;

import java.io.IOException;

import org.apache.hadoop.fs.Path;

```



```

import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.InputSplit;
import org.apache.hadoop.mapreduce.RecordReader;
import org.apache.hadoop.mapreduce.TaskAttemptContext;
import org.apache.hadoop.mapreduce.lib.input.CombineFileSplit;
;
import org.apache.hadoop.mapreduce.lib.input.FileSplit;
import org.apache.hadoop.mapreduce.lib.input.LineRecordReader;

public class CombineSmallfileRecordReader extends RecordReader<LongWritable, Text> {

    private CombineFileSplit combineFileSplit;//分片
    private LineRecordReader lineRecordReader = new LineRecordReader();//读文件
    private Path[] paths;//路径
    private int totalLength;//总长度
    private int currentIndex;
    private float currentProgress = 0;
    private LongWritable currentKey;
    private Text currentValue = new Text();

    public CombineSmallfileRecordReader(CombineFileSplit combineFileSplit, TaskAttemptContext context, Integer index) throws IOException {
        super();
        this.combineFileSplit = combineFileSplit;
        // 当前要处理的小文件Block在CombineFileSplit中的索引
        this.currentIndex = index;
    }
    /**
     * 循环初始化，此方法会被循环调用。
     */
    @Override
    public void initialize(InputSplit split, TaskAttemptContext context) throws IOException, InterruptedException {
        //分片
        this.combineFileSplit = (CombineFileSplit) split;
        //初始化每一片的io
        FileSplit fileSplit = new FileSplit(combineFileSplit.getPath(currentIndex), combineFileSplit.getOffset(currentIndex), combineFileSplit.getLength(currentIndex), combineFileSplit.getLocations());
        lineRecordReader.initialize(fileSplit, context);
        //获取所有路径
        this.paths = combineFileSplit.getPaths();
        //几个文件块，如果文件中包含大于64M的文件，算两个文件块，如241M的文件还会被视为4个文件块
        totalLength = paths.length;
        //设置当前处理的文件的文件名
        context.getConfiguration().set("map.input.file.name", combineFileSplit.getPath(currentIndex).getName());
    }
    @Override
    public boolean nextKeyValue() throws IOException, InterruptedException {
        //如果当前文件的索引大于0，并且没有超过所有文件的总长度
        if (currentIndex >= 0 && currentIndex < totalLength) {
            //读一行，记录到lineRecordReader的成员变量key、value中
            return lineRecordReader.nextKeyValue();
        } else {
            return false;
        }
    }
    @Override
    public LongWritable getCurrentKey() throws IOException, InterruptedException {
        //返回之前读取的key
        currentKey = lineRecordReader.getCurrentKey();
        return currentKey;
    }

    @Override
    public Text getCurrentValue() throws IOException, InterruptedException {
        //返回之前读取的value
        currentValue = lineRecordReader.getCurrentValue();
        return currentValue;
    }

    /**
     * 执行流程
     */

```

```
@Override
public float getProgress() throws IOException {
    if (currentIndex >= 0 && currentIndex < totalLength) {
        currentProgress = (float) currentIndex / totalLength;
        return currentProgress;
    }
    return currentProgress;
}
/**
 * 关闭流
 */
@Override
public void close() throws IOException {
    lineRecordReader.close();
}
}
```